

CAUSAL SCENE FLOW FOR INTERACTION: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study 3D perception for manipulation. The motivating bet is: Infer which scene-flow components are caused by robot action versus passive motion.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For 3D perception for manipulation, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2026a; 2025; 2019a; 2026b; 2024; 2019b). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.575	0.425
observed_only_wam	hidden_mode_shift	0.185	0.815
augmentation_as_data_wam	nominal	0.600	0.400
augmentation_as_data_wam	hidden_mode_shift	0.287	0.713
scalar_uncertainty_wam	nominal	0.667	0.333
scalar_uncertainty_wam	hidden_mode_shift	0.325	0.675
proposed_branch_mechanism	nominal	0.740	0.260
proposed_branch_mechanism	hidden_mode_shift	0.562	0.438

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Causal Scene Flow for Interaction argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Natural teaching for humanoid robot via human-in-the-loop scene-motion cross-modal perception, 2019a. <https://doi.org/10.1108/ir-06-2018-0118>.
- Unsupervised learning of a hierarchical spiking neural network for optical flow estimation: From events to global motion perception, 2019b. <https://doi.org/10.1109/tpami.2019.2903179>.
- Duel: Depth visual ego-motion learning for autonomous robot obstacle avoidance, 2024. <https://doi.org/10.1177/02783649231210325>.
- Hribench: Benchmarking vision-language models for real-time human perception in human-robot interaction, 2025. <http://arxiv.org/abs/2506.20566v1>.
- Aerial manipulation with contact-aware onboard perception and hybrid control, 2026a. <http://arxiv.org/abs/2602.08251v2>.
- Sensing the action: Rethinking sensor modalities and multi-modal fusion in vision-language-action models for robotic manipulation, 2026b. <https://doi.org/10.3390/s26113541>.